Bringing Clarity and Safety to the Miracle of Life!

By- Divyanshu Singh
(B.S Data Science & Applications, IIT Madras)

PURPOSE

Every expectant parent holds a treasure of hope and anticipation for the arrival of their baby. Amidst this joy, uncertainty can cast shadows – particularly when it comes to the health and well-being of both the mother and the developing baby. This project, focused on fetal head segmentation, is more than just a technical endeavor; it's an emotional journey that's driven by a powerful desire to provide clarity, assurance, and safety during a pivotal phase of life.

Imagine the heartbeats that race as parents-to-be enter ultrasound rooms, eagerly waiting to see glimpses of their child's first images. Now, envision the relief and comfort that flood their hearts when they are presented with a clear, detailed, and accurate visual representation of their baby's head – a crucial region for gauging development and health. This project is fueled by the empathy for those parents, standing on the threshold between their dreams and the reality of pregnancy.

The emotional impact of this work is deeply profound. It's about allowing medical professionals to make informed decisions that can influence both the delivery process and the well-being of the mother and child. It's about giving parents a tangible connection to the life growing within them, painting a picture of hope, progress, and care. With each pixel that's accurately segmented, you're playing a role in alleviating the anxiety that often comes with medical uncertainty.

Beyond the technical aspects of deep learning and image segmentation, this project holds the potential to alleviate fears, strengthen bonds, and shape brighter futures. It's a pursuit driven not just by data, but by the universal language of empathy, understanding, and the shared experience of parenthood. Through this endeavor, you are contributing to a journey of compassion, ensuring that every parent feels supported and connected during one of the most significant and transformative times of their lives.

Project Approach and Achievements

In this project, I embarked on a comprehensive journey to address the task of fetal head segmentation using Convolutional Neural Networks (CNNs) implemented in PyTorch. The goal was to develop an automated model capable of accurately outlining the boundary of fetal heads in ultrasound images, thereby aiding medical professionals and expectant parents.

Dataset Preprocessing:

I began by downloading and extracting the provided training dataset, which contained ultrasound images along with corresponding annotations in the form of binary masks. I split the dataset into a training set (700 images) and a validation set (300 images) to facilitate model training and evaluation.

Residual UNet Model vs. Plain Encoder-Decoder:

In my pursuit of achieving accurate fetal head segmentation, I ventured into two distinct architectural paradigms: the Residual UNet model and the plain Encoder-Decoder model.

- **Residual UNet Model:** Utilizing the Residual Network (ResNet) architecture within the UNet framework proved to be a highly effective strategy. This model demonstrated remarkable segmentation accuracy, achieving an impressive Dice score of 91%. The inclusion of residual connections in the UNet architecture allowed for better feature propagation, enhancing the model's ability to capture intricate details in ultrasound images.
- **Plain Encoder-Decoder Model:** On the other hand, the plain Encoder-Decoder model exhibited a Dice score of 74%. Despite its simplicity, this architecture failed to match the segmentation accuracy of the Residual UNet. The absence of residual connections hindered the model's capacity to effectively capture nuanced features, leading to a lower performance level.

Training and Validation Analysis:

Throughout the training process, I meticulously tracked the training and validation losses to monitor model convergence and generalization. By visualizing these metrics over epochs, I gained insights into the models' learning trajectories and their adaptability to the dataset.

Deployment on Test Set:

After honing the Residual UNet model to achieve the impressive Dice score of 91%, I deployed the model on the test set, comprising 335 ultrasound images. The deployment successfully showcased the model's proficiency in real-world scenarios. The generated segmentations effectively highlighted fetal head boundaries, validating the model's utility in practical medical applications.

Conclusion:

In conclusion, my project journey encompassed rigorous experimentation, architectural exploration, and meticulous evaluation. The Residual UNet model's triumphant performance underscored the significance of architectural innovations in enhancing segmentation accuracy.

This project not only achieved remarkable results in fetal head segmentation but also underscored the potential of CNNs to revolutionize medical image analysis and contribute to advanced healthcare practices.